

Double-click (or enter) to edit

```
!pip install numpy seaborn pandas matplotlib scikit-learn
```

```
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (1.25.2)
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (0.13.1)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (2.0.3)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.2.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.2.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (4.53.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (24.1)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (3.1.2)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.11.4)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.5.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)
```

```
import numpy as np
import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.linear_model import SGDClassifier, LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import accuracy_score
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier

from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
from sklearn.feature_selection import f_regression
from sklearn.metrics import mean_squared_error
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report

import warnings
warnings.filterwarnings('ignore')

df = pd.read_csv('Flight delay prediction.csv')
```

Data Preprocessing

```
df.columns

Index(['id', 'Airline', 'Flight', 'AirportFrom', 'AirportTo', 'DayOfWeek',
      'Time', 'Length', 'Delay', 'weather'],
      dtype='object')
```

```
df['Delay'].value_counts()

Delay
0    299119
1     240264
Name: count, dtype: int64
```

```
df['weather'].value_counts()

weather
Perfect      299119
Imperfect    240264
Name: count, dtype: int64
```

```
df['Airline'].value_counts()

Airline
WN      94097
```

DL 60940  
OO 50254  
AA 45656  
MQ 36605  
US 34500  
XE 31126  
EV 27983  
UA 27619  
CO 21118  
FL 20827  
9E 20686  
B6 18112  
YV 13725  
OH 12630  
AS 11471  
F9 6456  
HA 5578  
Name: count, dtype: int64

df.info()

<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 539383 entries, 0 to 539382  
Data columns (total 10 columns):  
# Column Non-Null Count Dtype  
0 id 539383 non-null int64  
1 Airline 539383 non-null object  
2 Flight 539383 non-null int64  
3 AirportFrom 539383 non-null object  
4 AirportTo 539383 non-null object  
5 DayOfWeek 539383 non-null int64  
6 Time 539383 non-null int64  
7 Length 539383 non-null int64  
8 Delay 539383 non-null int64  
9 weather 539383 non-null object  
dtypes: int64(6), object(4)  
memory usage: 41.2+ MB

df.isnull().sum()

id 0  
Airline 0  
Flight 0  
AirportFrom 0  
AirportTo 0  
DayOfWeek 0  
Time 0  
Length 0  
Delay 0  
weather 0  
dtype: int64

df.describe()

	id	Flight	DayOfWeek	Time	Length	Delay
count	539383.000000	539383.000000	539383.000000	539383.000000	539383.000000	539383.000000
mean	269692.000000	2427.928630	3.929668	802.728963	132.202007	0.445442
std	155706.604461	2067.429837	1.914664	278.045911	70.117016	0.497015
min	1.000000	1.000000	1.000000	10.000000	0.000000	0.000000
25%	134846.500000	712.000000	2.000000	565.000000	81.000000	0.000000
50%	269692.000000	1809.000000	4.000000	795.000000	115.000000	0.000000
75%	404537.500000	3745.000000	5.000000	1035.000000	162.000000	1.000000

df.head()

	id	Airline	Flight	AirportFrom	AirportTo	DayOfWeek	Time	Length	Delay	weather
0	1	CO	269	SFO	IAH	3	15	205	1	Imperfect
1	2	US	1558	PHX	CLT	3	15	222	1	Imperfect
2	3	AA	2400	LAX	DFW	3	20	165	1	Imperfect
3	4	AA	2466	SFO	DFW	3	20	195	1	Imperfect

```
# List of categorical columns to be label encoded
categorical_columns = ['Airline', 'AirportFrom', 'AirportTo','weather']

# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Iterate over each categorical column and transform it
for col in categorical_columns:
    df[col] = label_encoder.fit_transform(df[col])
```

## Feature selection

```
# Separate features (X) and target variable (y)
X = df.drop(columns=['Delay','id']) # Exclude the target column from features
y = df['Delay']

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize SelectKBest with ANOVA F-value scoring
k_best = SelectKBest(score_func=f_regression, k=8)
# Fit SelectKBest to training data
X_train_kbest = k_best.fit_transform(X_train, y_train)

# Get indices of selected features
selected_indices = k_best.get_support(indices=True)

# Get names of selected features
selected_features = X.columns[selected_indices]

# Evaluate the selected features
# Example: Fit a model and evaluate its performance
model = LinearRegression()
model.fit(X_train_kbest, y_train)
y_pred = model.predict(k_best.transform(X_test))
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error (Using SelectKBest):", mse)
print("Selected features:", selected_features)
```

```
➦ Mean Squared Error (Using SelectKBest): 1.0849515342250246e-29
   Selected features: Index(['Airline', 'Flight', 'AirportFrom', 'AirportTo', 'DayOfWeek', 'Time',
                           'Length', 'weather'],
                           dtype='object')
```

## Train Test Split

```
X = df[['Airline', 'Flight', 'AirportFrom', 'AirportTo', 'DayOfWeek', 'Time', 'Length', 'weather']].values
y = df['Delay'].values
```

```
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size = 0.2, random_state=5)
```

## Model implementation


### Stochastic Gradient Descent (SGD)

```
sgd = SGDClassifier(max_iter=5, tol=None)
sgd.fit(X_train, y_train)
Y_pred = sgd.predict(X_test)
sgd.score(X_train, y_train)
acc_sgd = round(sgd.score(X_train, y_train) * 100, 2)
print(round(acc_sgd,2), "%")
```

```
➦ 47.05 %
```

```
# Calculate confusion matrix
conf_mat = confusion_matrix(y_test, Y_pred)

# Plot confusion matrix as a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(conf_mat, annot=True, fmt='d', cmap="Blues", xticklabels=['Not Delayed', 'Delayed'], yticklabels=['Not Delayed', 'Delayed'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
# Print classification report
print("Classification Report:")
print(classification_report(y_test, Y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.05	0.09	59768
1	0.46	1.00	0.63	48109
accuracy			0.47	107877
macro avg	0.72	0.52	0.36	107877
weighted avg	0.75	0.47	0.33	107877

## Random Forest Classifier

```
random_forest = RandomForestClassifier(n_estimators=100)
random_forest.fit(X_train, y_train)

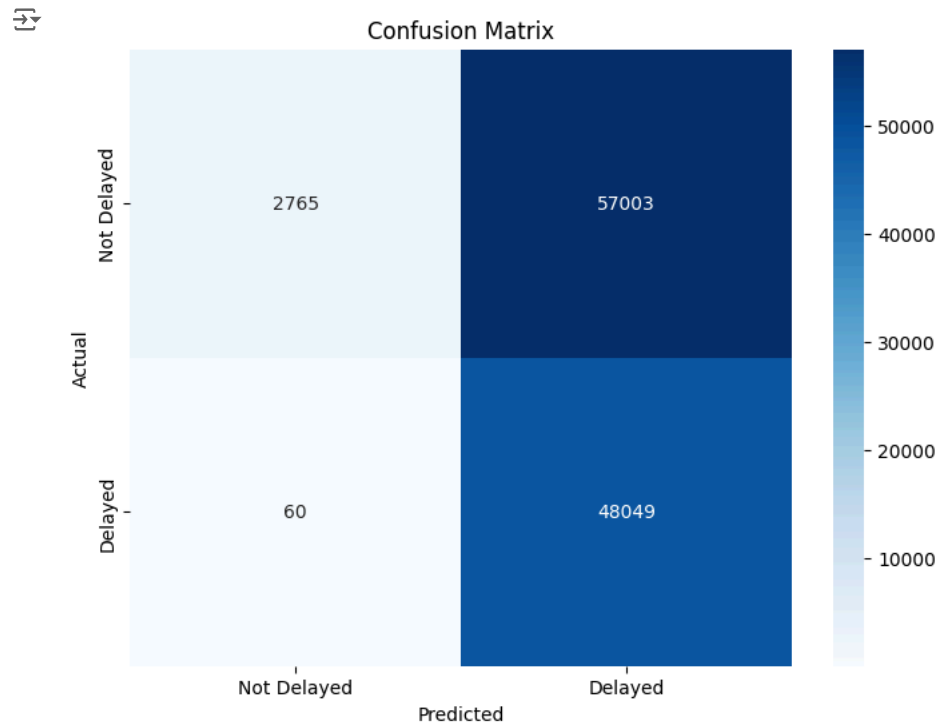
Y_prediction = random_forest.predict(X_test)

random_forest.score(X_train, y_train)
acc_random_forest = round(random_forest.score(X_train, y_train) * 100, 2)
print(round(acc_random_forest,2), "%")
```

100.0 %

```
# Calculate confusion matrix
conf_mat = confusion_matrix(y_test, Y_pred)

# Plot confusion matrix as a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(conf_mat, annot=True, fmt='d', cmap="Blues", xticklabels=['Not Delayed', 'Delayed'], yticklabels=['Not Delayed', 'Delayed'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
# Print classification report
print("Classification Report:")
print(classification_report(y_test, Y_pred))
```

Classification Report:					
	precision	recall	f1-score	support	
0	0.98	0.05	0.09	59768	
1	0.46	1.00	0.63	48109	
accuracy			0.47	107877	
macro avg	0.72	0.52	0.36	107877	
weighted avg	0.75	0.47	0.33	107877	

## Logistic Regression

```
logreg = LogisticRegression()
logreg.fit(X_train, y_train)


Y_pred = logreg.predict(X_test)

acc_log = round(logreg.score(X_train, y_train) * 100, 2)
print(round(acc_log,2), "%")
```

100.0 %

```
# Calculate confusion matrix
conf_mat = confusion_matrix(y_test, Y_pred)

# Plot confusion matrix as a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(conf_mat, annot=True, fmt='d', cmap="Blues", xticklabels=['Not Delayed', 'Delayed'], yticklabels=['Not Delayed', 'Delayed'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
# Print classification report
print("Classification Report:")
print(classification_report(y_test, Y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	59768
1	1.00	1.00	1.00	48109
accuracy			1.00	107877
macro avg	1.00	1.00	1.00	107877
weighted avg	1.00	1.00	1.00	107877

### K-Nearest Neighbor (KNN)


```
knn = KNeighborsClassifier(n_neighbors = 3)
knn.fit(X_train, y_train)

Y_pred = knn.predict(X_test)
acc_knn = round(knn.score(X_train, y_train) * 100, 2)
print(round(acc_knn,2), "%")
```

91.24 %

```
# Calculate confusion matrix
conf_mat = confusion_matrix(y_test, Y_pred)

# Plot confusion matrix as a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(conf_mat, annot=True, fmt='d', cmap="Blues", xticklabels=['Not Delayed', 'Delayed'], yticklabels=['Not Delayed', 'Delayed'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
# Print classification report
print("Classification Report:")
print(classification_report(y_test, Y_pred))
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.79       0.83       0.81     59768
     1       0.77       0.72       0.75     48109

 accuracy          0.78
 macro avg          0.78
weighted avg          0.78
```


## GaussianNB

```
nb= GaussianNB()
nb.fit(X_train, y_train)
nb_pred = nb.predict(X_test)
nb_accuracy = round(accuracy_score(y_test, nb_pred)* 100, 2)
print("Naive Bayes Accuracy:", nb_accuracy)
print(round(nb_accuracy,2), "%")
```

```
Naive Bayes Accuracy: 100.0
100.0 %
```

```
# Calculate confusion matrix
conf_mat = confusion_matrix(y_test, nb_pred)

# Plot confusion matrix as a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(conf_mat, annot=True, fmt='d', cmap="Blues", xticklabels=['Not Delayed', 'Delayed'], yticklabels=['Not Delayed', 'Delayed'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
# Print classification report
print("Classification Report:")
print(classification_report(y_test, nb_pred))
```

```
Classification Report:
              precision    recall  f1-score   support

     0       1.00        1.00        1.00     59768
     1       1.00        1.00        1.00     48109

 accuracy          1.00
 macro avg          1.00
weighted avg          1.00
```

## DecisionTreeClassifier

```
dt= DecisionTreeClassifier()
dt.fit(X_train, y_train)
dt_pred = dt.predict(X_test)
dt_accuracy = round(accuracy_score(y_test, dt_pred)* 100, 2)
print("Decision Tree Accuracy:", dt_accuracy)
```

```
# Calculate confusion matrix
conf_mat = confusion_matrix(y_test, dt_pred)
```


```
Decision Tree Accuracy: 100.0
```

```
# Plot confusion matrix as a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(conf_mat, annot=True, fmt='d', cmap="Blues", xticklabels=['Not Delayed', 'Delayed'], yticklabels=['Not Delayed', 'Delayed'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```





Confusion Matrix



```
# Print classification report
print("Classification Report:")
print(classification_report(y_test, dt_pred))
```



```
Classification Report:
              precision    recall  f1-score   support

     0       1.00        1.00        1.00     59768
     1       1.00        1.00        1.00     48109

 accuracy          1.00
 macro avg          1.00
weighted avg          1.00
```

## VOTING CLASSIFIER

```
from sklearn.ensemble import VotingClassifier

# Define the individual models
model1 = GaussianNB()
model2 = DecisionTreeClassifier()

# Define the ensemble model
model = VotingClassifier(estimators=[('nb', model1), ('dt', model2)], voting='hard')

# Fit the model to the training data
model.fit(X_train, y_train)

# Make predictions
model_pred = model.predict(X_test)

# Calculate accuracy
model_accuracy = round(accuracy_score(y_test, model_pred) * 100, 2)
print("Hybrid Model Accuracy:", model_accuracy,"%")

# Print classification report
print("Classification Report:")
print(classification_report(y_test, model_pred))
```



```
Hybrid Model Accuracy: 100.0 %
Classification Report:
              precision    recall  f1-score   support

     0       1.00        1.00        1.00     59768
     1       1.00        1.00        1.00     48109

 accuracy          1.00
 macro avg          1.00
weighted avg          1.00
```

## RNN

```
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers, models # Add this line to import the models submodule

# Assuming X_train, X_test, y_train, y_test are already prepared

# Define the RNN model
model_rnn = models.Sequential([ # This line should work now
    layers.SimpleRNN(32, input_shape=(X_train.shape[1], 1)),
    layers.Dense(1, activation='sigmoid')
])

# Compile the model
model_rnn.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Evaluate the model and calculate accuracy
evaluation_result = model_rnn.evaluate(X_test, y_test)
accuracy_rnn = round(evaluation_result[1] * 100, 2)
print("Accuracy of RNN model:", accuracy_rnn, "%")
```

3372/3372 [=====] - 7s 2ms/step - loss: 0.6839 - accuracy: 0.5540  
Accuracy of RNN model: 55.4 %

## LSTM

```
# Assuming X_train, X_test, y_train, y_test are already prepared
# Define the LSTM model
model_lstm = models.Sequential([
    layers.LSTM(32, input_shape=(X_train.shape[1], 1)),
    layers.Dense(1, activation='sigmoid')
])

# Compile the model
model_lstm.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Train the model
model_lstm.fit(X_train, y_train, epochs=1, batch_size=32, validation_data=(X_test, y_test))

# Evaluate the model
evaluation_result = model_lstm.evaluate(X_test, y_test)
accuracy_lstm = round(evaluation_result[1] * 100, 2)
print("Accuracy of LSTM model:", accuracy_lstm, "%")
```

13485/13485 [=====] - 82s 6ms/step - loss: 0.0121 - accuracy: 0.9972 - val\_loss: 3.6054e-05 - val\_accuracy: 1.0000  
3372/3372 [=====] - 9s 3ms/step - loss: 3.6054e-05 - accuracy: 1.0000  
Accuracy of LSTM model: 100.0 %

## CNN

```
# Reshape the input data for CNN
X_train_cnn = np.reshape(X_train, (X_train.shape[0], X_train.shape[1], 1))
X_test_cnn = np.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))

# Define the CNN model
model_cnn = models.Sequential([
    layers.Conv1D(32, 3, activation='relu', input_shape=(X_train_cnn.shape[1], 1)),
    layers.MaxPooling1D(2),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(1, activation='sigmoid')
])

# Compile the model
model_cnn.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Evaluate the model and calculate accuracy
evaluation_result = model_cnn.evaluate(X_test_cnn, y_test)
accuracy_cnn = round(evaluation_result[1] * 100, 2)
print("Accuracy of CNN model:", accuracy_cnn, "%")
```

3372/3372 [=====] - 5s 1ms/step - loss: 49.1335 - accuracy: 0.4798  
Accuracy of CNN model: 47.98 %

## LSTM with SGD [Hybrid]

```

from tensorflow.keras import layers, models
from tensorflow.keras.optimizers import SGD

# Assuming X_train, X_test, y_train, y_test are already prepared

# Define the LSTM model
model_lstm = models.Sequential([
    layers.LSTM(32, input_shape=(X_train.shape[1], 1)),
    layers.Dense(1, activation='sigmoid')
])

# Define the SGD optimizer without the decay argument
sgd = SGD(learning_rate=0.01, momentum=0.9, nesterov=True)

# Compile the model with SGD optimizer
model_lstm.compile(optimizer=sgd, loss='binary_crossentropy', metrics=['accuracy'])

# Train the model
model_lstm.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test, y_test))

# Evaluate the model
evaluation_result = model_lstm.evaluate(X_test, y_test)
accuracy_lstm = round(evaluation_result[1] * 100, 2)
print("Accuracy of LSTM model with SGD optimizer:", accuracy_lstm, "%")

```



```

Epoch 1/10
13485/13485 [=====] - 82s 6ms/step - loss: 0.0157 - accuracy: 0.9956 - val_loss: 2.9158e-04 - val_accuracy:
Epoch 2/10
13485/13485 [=====] - 75s 6ms/step - loss: 1.9211e-04 - accuracy: 1.0000 - val_loss: 1.3269e-04 - val_accuracy:
Epoch 3/10
13485/13485 [=====] - 75s 6ms/step - loss: 1.0474e-04 - accuracy: 1.0000 - val_loss: 8.4416e-05 - val_accuracy:
Epoch 4/10
13485/13485 [=====] - 79s 6ms/step - loss: 7.1461e-05 - accuracy: 1.0000 - val_loss: 6.1419e-05 - val_accuracy:
Epoch 5/10
13485/13485 [=====] - 76s 6ms/step - loss: 5.3958e-05 - accuracy: 1.0000 - val_loss: 4.8055e-05 - val_accuracy:
Epoch 6/10
13485/13485 [=====] - 76s 6ms/step - loss: 4.3200e-05 - accuracy: 1.0000 - val_loss: 3.9360e-05 - val_accuracy:
Epoch 7/10
13485/13485 [=====] - 73s 5ms/step - loss: 3.5936e-05 - accuracy: 1.0000 - val_loss: 3.3267e-05 - val_accuracy:
Epoch 8/10
13485/13485 [=====] - 74s 6ms/step - loss: 3.0711e-05 - accuracy: 1.0000 - val_loss: 2.8772e-05 - val_accuracy:
Epoch 9/10
13485/13485 [=====] - 71s 5ms/step - loss: 2.6781e-05 - accuracy: 1.0000 - val_loss: 2.5316e-05 - val_accuracy:
Epoch 10/10
13485/13485 [=====] - 72s 5ms/step - loss: 2.3720e-05 - accuracy: 1.0000 - val_loss: 2.2578e-05 - val_accuracy:

```